Valentino Trend Forecasting

A Data-Driven Forecast for SS 2027 and FW 2027

Independent Research Project

March 2026

Executive Summary

This report presents the findings of a machine learning project commissioned to forecast what Valentino's next two runway collections may look like. Using eight seasons of Valentino runway imagery and the Claude Vision AI, I built a structured dataset of 617 looks, identified visual trends across seasons, and applied two independent forecasting methods to predict the SS 2027 and FW 2027 collections. Each forecast was then used to generate a visualized look using DALL-E 3, OpenAI's image generation model. The result is a concrete, visual prediction of where Valentino's aesthetic is heading.

The Brief

The client engaged this project with a straightforward question: can data tell us what a luxury fashion house will do next? Valentino was selected as the subject given its distinct visual identity and its availability of documented runway imagery across multiple seasons. Rather than relying on editorial intuition alone, the goal was to build a reproducible, data-grounded approach to trend forecasting that could eventually be applied to any designer.

The project was scoped to cover eight consecutive seasons: four Spring/Summer and four Fall/Winter from 2023 through 2026, and to produce forward-looking predictions for both SS 2027 and FW 2027.

Approach

The project was built as a four-stage pipeline:

Stage 1) Image Analysis: Each of the 617 runway photographs was individually analyzed by Claude, Anthropic's multimodal large language model. Unlike a traditional computer vision classifier, which is trained to output a fixed set of labels, Claude is a transformer-based model that processes images and text through the same attention mechanism. The image is divided into a grid of patches, each encoded into a numerical vector and treated as a token, exactly like a word in a sentence. Claude then attends over both the image tokens and the text prompt simultaneously, allowing it to generate a nuanced, context-sensitive description rather than choosing from a predetermined list. We instructed Claude to return its output as structured JSON, specifying six attributes per image: dominant colors, silhouette type, key garments, fabrics, patterns, and overall aesthetic. This turned a folder of photographs into a structured spreadsheet that could be analyzed statistically.

Stage 2) Trend Analysis: With the data in hand, we computed frequency distributions for each attribute across all eight seasons, tracking how the presence of specific colors, silhouettes, and aesthetics shifted from season to season.

Stage 3) Forecasting: Two independent methods were applied. The first, a statistical approach, uses the frequency data itself to project trends forward, either by weighting recent seasons more heavily, or by fitting a line to each trend and projecting where it is heading. The second method feeds the trend summary back to Claude and asks it to reason through what the next season will look like. This is a form of in-context learning: rather than fine-tuning a model specifically on fashion data, we supply the trend evidence directly in the prompt and allow the model to reason over it alongside its parametric knowledge. The fashion understanding is baked into its weights during training on billions of documents. The result is a synthesis of what the data shows and what the model already knows about how luxury houses evolve. Claude Sonnet was used for this step rather than the lighter Claude Haiku model used for image tagging. This was a deliberate choice, since forecasting requires multi-step reasoning and the ability to hold seasonal distinctions in mind, which benefits from a more capable model.

Stage 4) Image Generation: Both forecasts were converted into detailed visual descriptions and passed to DALL-E 3, OpenAI's image generation model, producing a runway-style photograph of each predicted look.

What the Data Shows

SS 2027 Trend Analysis

The chart below summarizes the statistical forecast for SS 2027 across colors, silhouette, and aesthetic. Two methods are shown: Method 1 (Weighted Recency) emphasizes the most recent SS seasons, while Method 2 (Linear Trend) projects the directional trajectory of each value forward.

Figure 1. Statistical forecast summary for Valentino SS 2027. Each bar represents the predicted relative share of that attribute in the upcoming collection.

Both methods agree that black remains the dominant color, with gold and cream as strong supporting tones. Silhouette predictions diverge slightly — weighted recency favors tailored and flowing shapes,

while the linear trend points more strongly toward tailored and structured. Both converge on haute couture and avant-garde as the defining aesthetic registers.

FW 2027 Trend Analysis

The FW forecast tells a somewhat different story, reflecting the distinct visual language that Valentino typically employs in its winter collections.

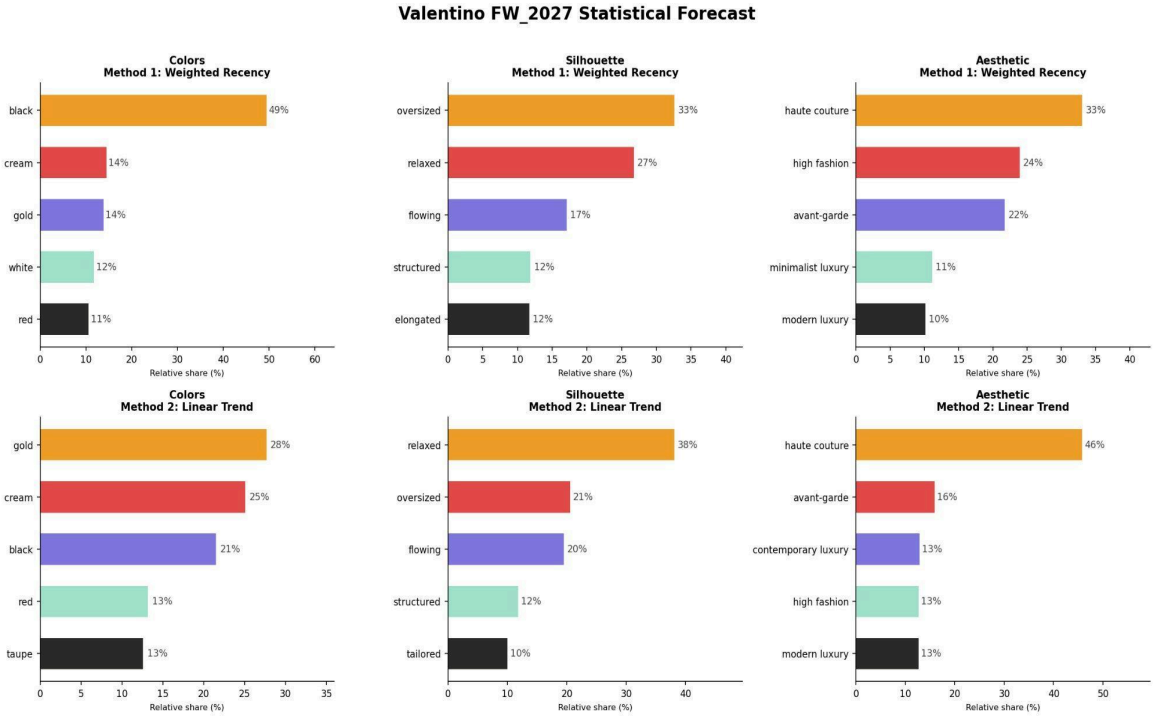

Figure 2. Statistical forecast summary for Valentino FW 2027. Note the shift toward oversized and relaxed silhouettes compared to the SS forecast, consistent with FW-specific patterns in the historical data.

For FW 2027, black is again dominant but gold rises notably in the linear trend method — suggesting it is gaining share faster in FW collections than SS. The silhouette picture shifts meaningfully: oversized, relaxed, and flowing structures come forward, replacing the tailored precision predicted for SS. This seasonal distinction is exactly what we hoped the model would detect, and it does so clearly.

The Predicted Looks

SS 2027: Claude API Forecast

Based on the trend summary, Claude's reasoning model predicted a collection defined by structured, military-inspired tailoring in a black and gold palette, with elongated silhouettes and haute couture embellishment. The generated image below visualizes this prediction.

Figure 3. DALL-E 3 generated look based on Claude's SS 2027 forecast. A structured black military-style jacket with gold embroidery and a fitted midi silhouette, styled with gold heeled boots.

SS 2027: Statistical Forecast

The statistical method, working from frequency projections alone, arrived at a complementary but subtly different prediction, emphasizing the elongated, androgynous quality that the linear trend model identified as a rising signal in SS collections.

Figure 4. DALL-E 3 generated look based on the statistical SS 2027 forecast. An elongated black coat with gold piping over a flowing cream draped skirt, reflecting the androgynous and contemporary luxury signals surfaced by the linear trend projection.

FW 2027: Claude API Forecast

For the winter collection, Claude predicted a return to the oversized, enveloping silhouettes that defined Valentino's earlier FW seasons. This is reinterpreted through a more refined, contemporary luxury lens, with black and taupe tones and heavy felted fabrics.

Figure 5. DALL-E 3 generated look based on Claude's FW 2027 forecast. Dramatically oversized black and taupe wool coats with clean geometric proportions, evoking modern elegance and architectural volume.

FW 2027: Statistical Forecast

The statistical model pushed further into drama and volume, reflecting the strong avant-garde and haute couture signals in the FW historical data. The resulting image captures a more theatrical interpretation of the same underlying trend.

Figure 6. DALL-E 3 generated look based on the statistical FW 2027 forecast. A dramatic oversized hooded black coat with embroidered sleeves over a cream draped inner layer, with gold boots — reflecting the haute couture and avant-garde signals dominant in FW historical data.

Comparing the Two Methods

Both approaches produced credible, visually coherent predictions, and they agreed on the fundamentals: black dominates, gold is ascending, and the haute couture register remains central to Valentino's identity. Where they diverge is instructive.

The statistical method tends to surface signals that are growing even when they are not yet dominant. In the SS forecast, it picked up on androgynous silhouettes and modern luxury aesthetics before they reached the top of the frequency rankings. These are early-stage trends that weighted recency averaging would not yet flag. In the FW forecast, it leaned further into theatrical avant-garde volume, amplifying signals that the Claude model tempered with contextual judgment.

Claude's reasoning-based forecast, by contrast, produced more cohesive and editorially grounded predictions. It correctly separated SS and FW seasonal patterns and synthesized the trend data with broader fashion knowledge to produce looks that feel wearable and directionally coherent. The trade-off is that it may be more conservative and therefore less likely to surface emerging signals before they become obvious. There is also an important technical caveat: LLM outputs are stochastic. Claude does not produce the same answer every time; it samples from a probability distribution over possible next tokens, meaning the forecast represents one plausible synthesis rather than a deterministic conclusion. Running the same prompt multiple times would produce a family of related but distinct predictions, all consistent with the data.

Used together, the two methods offer complementary lenses: the statistical model as an early-warning system for emerging directions, and the Claude model as a contextualizing interpreter.

Limitations and Next Steps

Eight seasons is a meaningful but relatively small dataset for statistical forecasting. Trend projections should be read as directional signals rather than precise predictions, and the linear method in particular is sensitive to the volatility of individual season data points.

The free-text nature of Claude's image analysis also introduces some inconsistency in vocabulary. Two looks with the same silhouette might be tagged slightly differently across the dataset. A future version of this project could standardize the taxonomy using a controlled vocabulary, improving the statistical reliability of the frequency data.

The most natural next step would be to validate the methodology retrospectively: use seasons 1 through 6 to predict season 7, then compare the prediction against what was actually shown. This would give a concrete measure of forecast accuracy and allow the approach to be calibrated before being used for forward-looking decisions.